

SHYAM KUMAR

AI Governance & AI Red-Teaming | Cybersecurity | NIST AI RMF · OWASP LLM Top 10 · MITRE ATLAS

Chennai, Tamil Nadu, India | shyamvasansathiskumar@gmail.com | +91 93635 92140 | [LinkedIn](#) | [GitHub](#)

SUMMARY

CSE undergraduate focused on AI security and AI governance — red-teaming local LLMs with garak, mapping the OWASP LLM Top 10 to MITRE ATLAS technique by technique, and writing risk assessments against the NIST AI RMF. Building both a technical and a governance track in public, documented across three GitHub repositories. Early in the field and building deliberately toward AI red-teaming / AI governance roles.

PROJECTS

redteam-log | Personal project — AI red-teaming lab notebook — Aug 2026 – Present

Hands-on AI security practice using garak and PyRIT against local LLMs (Ollama)

- Mapped OWASP LLM Top 10 (2025) vulnerability classes to their corresponding MITRE ATLAS techniques, starting with LLM01 (Prompt Injection → AML.T0051) and LLM02 (Sensitive Information Disclosure → AML.T0024), each with a full worked technical example
- Set up a reproducible red-team environment (garak + PyRIT + Ollama) with a documented install runbook and troubleshooting log
- Built structured templates for scan logs and CTF write-ups that force a “governance translation” step — turning a raw technical finding into a risk statement a non-technical stakeholder can act on

ai-governance-portfolio | Personal project — AI governance writing & risk assessment — Aug 2026 – Present

Applied writing kept deliberately separate from the technical lab notebook above

- Authored a full 9-section AI risk assessment template aligned to the NIST AI RMF 1.0 (system description, regulatory context, trustworthiness characteristics, risk register, RMF function mapping, limitations)
- Wrote a one-page AI use and governance policy template covering risk tiering, roles and responsibilities, pre-deployment requirements, and incident reporting

cybersecurity-fundamentals | Personal project — networking, web security, cryptography, CTF fundamentals — Sep 2026 – Present

The base layer underneath the two AI-specific repos above

- Logged a hand-verified subnetting exercise (/24 → eight /27s) quantifying the exact address cost of segmentation and its ARP-spoofing blast-radius implications
- Logged a reproducible AES ECB-vs-CBC demonstration (Python) showing 12,288 identical plaintext blocks collapse to 5 ciphertext blocks under ECB — the same structural leak behind Adobe's 2013 breach

LEADERSHIP & EXPERIENCE

Head, Tech & Skills Club | ASU Pathways — Sep 2026 – Present

- Founding head of a newly launched student club focused on technical skill-building; building out club structure and programming ahead of its first kickoff event

General Secretary | Velammal Bodhi Campus (Student Council) — 2023 – 2025

- Hosted school-wide events and secured sponsorships for student initiatives as part of the elected student leadership team

EDUCATION

Bachelor of Engineering, Computer Science and Engineering | Rajalakshmi Engineering College — Aug 2025 – May 2029

SGPA: 7.5

High School Diploma, Mathematics and Biology (PCMB) | Velammal Bodhi Campus — Feb 2023 – May 2025

Grade: 75% (First Class)

SKILLS

AI Security / Governance: AI Red-Teaming, AI Governance, Prompt Injection, NIST AI RMF, OWASP LLM Top 10, MITRE ATLAS, garak, PyRIT

Programming: Python, C, Data Structures, Algorithms, Git/GitHub

General: Cybersecurity fundamentals, technical writing, public speaking / event hosting

CERTIFICATIONS

- Anthropic (10+ issued, Sep 2026): AI Fluency: Framework & Foundations, Building with the Claude API, Model Context Protocol, Agents and Workflows, Prompt Engineering Techniques, and others
- NIST AI Risk Management Framework — in progress
- Microsoft AI Red Teaming 101 — in progress